

HOMEWORK 7

Elena Gómez

February 28, 2023

PROBLEM 0: Homework checklist

I worked alone.

PROBLEM 1:

(Adapted from Nocedal and Wright 3.1, typos are my own) Backtracking line search is a simple method that attempts to balance sufficient decrease with a long step length by testing a sequence of points:

$$\alpha = 1, 1/2, 1/4, \dots$$

and accepting the first point that satisfies sufficient decrease (which is the condition that $L(\alpha) \leq c_1 \alpha L'(0)$ for the line search function $L(\alpha)$ we discussed in class).

1. Implement a backtracking line search routine to select the step length and add this to our simple gradient descent code from the lectures. (Note, you are free to implement gradient descent using your own code too, adapting our code is only suggested for convenience.)

The following function *gradient_descendent_1* is the grandient descendent code with the changes neccesary for the backtracking line search. In this case, the descendent direction is $p = -g$.

In [185]:

```

1  using Printf, LinearAlgebra
2
3  function gradient_descent_1(fgh,x0;
4      maxiter=10000,tol=1.0e-8,quiet=false,histx=[],hista=[],gamma=0.01)
5
6      x = copy(x0)
7      n = length(x)
8
9      hist = zeros(2,maxiter)
10     savehistx = eltype(histx) == Vector{Float64} ? true : false
11     savehista = eltype(hista) == Float64 ? true : false
12     alpha_list=[]
13
14     f = Inf
15     normg = Inf
16     lastiter = 0
17     g = Vector{Float64}()
18     h = Matrix{Float64}(undef, 0, 0)
19
20
21     if !quiet
22         @printf(" %6s %9s %9s %9s\n", "iter",
23             "val", "normg", "fdiff");
24     end
25
26     for iter=1:maxiter
27         if savehistx
28             push!(histx, x)
29         end
30
31         if iter>1
32             p=-g; #direction
33
34             alpha=1.0;
35             while(fgh(x+alpha*p)[1] > f + gamma*alpha*p'*g)
36                 alpha/=2;
37             end
38
39             x=x+alpha*p;
40
41             push!(alpha_list,alpha);
42
43             if savehista
44                 push!(hista,alpha);
45             end
46         end
47
48         flast = f
49         f,g,h = fgh(x)
50         normg = norm(g,Inf)
51
52         fdiff = flast - f
53
54         if !quiet
55             @printf(" %6i %9.2e %9.2e %9.2e\n",
56                 iter, f, normg, fdiff)
57         end
58
59         hist[:,iter] = [f; normg]

```

```
60         lastiter = iter
61
62         if normg <= tol
63             break
64         end
65         if !isfinite(normg)
66             break
67         end
68     end
69
70     if lastiter < maxiter
71         hist = hist[:,1:lastiter]
72     end
73
74     if normg > tol
75         @warn "Did not converge"
76     end
77
78     return x,f,g,hist,alpha_list,histx
79 end
80 Out[185]:
```

gradient_descent_1 (generic function with 1 method)

The following code are auxiliar functions that we also use, as well as, the code we will use to plot the contour plots.

In [186]:

```

1 using Optim, OptimTestProblems
2 OptimTestProblems.UnconstrainedProblems.examples["Rosenbrock"]
3
4
5
6
7 using Plots
8 ezcontour(x, y, f) = begin
9     X = repeat(x', length(y), 1)
10    Y = repeat(y, 1, length(x))
11    # Evaluate each f(x, y)
12    Z = map((x,y) -> f([x,y]), X, Y)
13    plot(x, y, Z, st=:contour)
14 end
15
16
17
18
19
20 # These codes turn f and g into one function ...
21 function opt_combine(x, f, g!,h!)
22     g = Array{Float64,1}(undef,length(x))
23     h = Matrix{Float64}(undef, length(x), length(x));
24     g!(g,x)
25     h!(h,x)
26     return (f(x), g,h)
27 end
28
29 function opt_problem(p::OptimTestProblems.UnconstrainedProblems.OptimizationProblem)
30     return x -> opt_combine(x, p.f, p.g!,p.h!)
31 end
32
33 # this just makes it easy to use
34 opt_problem(s::AbstractString) = opt_problem(
35     OptimTestProblems.UnconstrainedProblems.examples[s])
36

```

Out[186]:

opt_problem (generic function with 2 methods)

2. Prepare a plot of the step lengths (values of α_k that were accepted) when we run gradient descent with this line search on the Rosenbrock function starting from the point (1.2, 1.2) and also the point (-1.2, 1).

Firstly, we compute the gradient descent starting from the point (1.2, 1.2)

In [187]:

```

1 opt_fun = MultivariateProblems.UnconstrainedProblems.examples["Rosenbrock"];
2 fgh = opt_problem(opt_fun);
3 histx = Vector{Vector{Float64}}();
4 hista = Vector{Float64}();
5
6
7 x,fx,gx,hist,alpha_list,histx1 = gradient_descent_1(fgh, [1.2,1.2];
8 maxiter=16000, tol=1.0e-8,quiet=false,
9 histx=histx, hista=hista, gamma=0.01);
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28 1.19e-02 2.05e-01 3.12e-06
29 1.18e-02 7.47e-02 3.68e-05
30 1.18e-02 2.77e-01 1.33e-05
31 1.18e-02 8.56e-02 5.64e-05
32 1.18e-02 1.61e-01 1.62e-05
33 1.18e-02 2.70e-01 3.90e-06
34 1.17e-02 9.76e-02 3.47e-05
35 1.17e-02 3.37e-01 1.64e-05
36 1.16e-02 1.02e-01 5.26e-05
37 1.16e-02 2.06e-01 1.72e-05
38 1.16e-02 1.86e-01 4.77e-06
39 1.16e-02 7.09e-02 3.23e-05
40 1.16e-02 2.48e-01 2.06e-05
41 1.15e-02 8.02e-02 4.79e-05
42 1.15e-02 1.52e-01 1.86e-05
43 1.15e-02 2.44e-01 6.02e-06
44 1.15e-02 9.48e-02 2.91e-05
45 1.14e-02 3.01e-01 2.54e-05
46 1.14e-02 9.85e-02 4.24e-05
47 1.14e-02 3.71e-01 3.81e-06

```


This method does not converge

Now, let's plot the values of α

In [188]:

```
1 plot(alpha_list)
```

Out[188]:




As we can see, this method takes too many iterations. In the plot above, we can see the values of α

Let's see the contour plot

In [190]:

```
1 ezcontour(-1.5:0.05:1.5, -1.5:0.05:1.5,  
2 MultivariateProblems.UnconstrainedProblems.examples["Rosenbrock"].f)  
3 #histx1 = histx1[1:3]  
4 plot!(map(first,histx1),map(x->x[2], histx1), linewidth=2)  
5
```

Out[190]:



Now, we compute the function starting from the point $(-1.2, 1)$

In [191]:

```
1 x,fx,gx,hist,alpha_list,histx1 = gradient_descent_1(fgh, [-1.2,1];
2 maxiter=16000, tol=1.0e-8,quiet=false,
3 histx=histx, hista=hista, gamma=0.01);
```


103	2.85e-03	8.60e-02	7.43e-06
104	2.85e-03	8.64e-02	6.47e-06
105	2.84e-03	7.04e-02	5.86e-06
106	2.84e-03	6.30e-02	5.47e-06
107	2.83e-03	9.83e-02	3.33e-06
108	2.82e-03	1.05e-01	7.53e-06
109	2.82e-03	7.82e-02	6.52e-06
110	2.81e-03	7.49e-02	5.87e-06
111	2.81e-03	6.55e-02	5.46e-06
112	2.80e-03	1.30e-01	2.77e-06
113	2.80e-03	8.84e-02	7.67e-06
114	2.79e-03	9.06e-02	6.58e-06
115	2.78e-03	7.19e-02	5.90e-06
116	2.78e-03	6.58e-02	5.46e-06
117	2.78e-03	1.02e-01	2.12e-06
118	2.77e-03	1.11e-01	7.84e-06
119	2.76e-03	8.05e-02	6.68e-06
120	2.76e-03	7.90e-02	5.94e-06
121	2.75e-03	6.68e-02	5.48e-06
122	2.75e-03	1.39e-01	1.35e-06

Now, let's plot the values of α

In [192]:

```
1 plot(alpha_list)
```

Out[192]:



Again, the method takes too many iterations. In the plot above we can see the values of α .

Let's see the contour plot


In [194]:

```

1 ezcontour(-1.5:0.05:1.5, -1.5:0.05:1.5,
2 MultivariateProblems.UnconstrainedProblems.examples["Rosenbrock"].f)
3 #histx1 = histx1[1:3]
4 plot!(map(first,histx1),map(x->x[2], histx1), linewidth=2)

```

Out[194]:



3. Implement Newton's method using line search as well. (Don't worry about singularity in the Hessian or any of the real issues that come up with Newton – just use the search direction

$$H(x_k)p = -g_k.$$

Let's compute the Newton's method using the linear search. The following function *newton* is the code with the changes necessary for the backtracking line search. The descent direction now is $H(x_k)p = -g_k$.

In [195]:

```

1  using Printf, LinearAlgebra
2
3  function newton(fgh,x0;
4      maxiter=10000,tol=1.0e-8,quiet=false,histx=[],hista=[],gamma=0.01)
5
6      x = copy(x0)
7      n = length(x)
8
9      hist = zeros(2,maxiter)
10     savehistx = eltype(histx) == Vector{Float64} ? true : false
11     savehista = eltype(hista) == Float64 ? true : false
12     alpha_list=[]
13
14
15     f = Inf
16     normg = Inf
17     lastiter = 0
18     g = Vector{Float64}()
19     h = Matrix{Float64}(undef, 0, 0)
20
21
22     if !quiet
23         @printf(" %6s %9s %9s %9s\n", "iter",
24             "val", "normg", "fdiff");
25     end
26
27     for iter=1:maxiter
28         if savehistx
29             push!(histx, x)
30         end
31
32         if iter>1
33             p=-h\g; #direction
34
35             alpha=1.0;
36             while(fgh(x+alpha*p)[1] > f + gamma*alpha*p'*g)
37                 alpha/=2;
38             end
39
40             x=x+alpha*p;
41
42             push!(alpha_list,alpha)
43
44             if savehista
45                 push!(hista,alpha);
46             end
47         end
48
49         flast = f
50         f,g,h = fgh(x)
51         normg = norm(g,Inf)
52
53         fdiff = flast - f
54
55         if !quiet
56             @printf(" %6i %9.2e %9.2e %9.2e\n",
57                 iter, f, normg, fdiff)
58         end
59

```

```

60     hist[:,iter] = [f; normg]
61     lastiter = iter
62
63     if normg <= tol
64         break
65     end
66     if !isfinite(normg)
67         break
68     end
69 end
70
71 if lastiter < maxiter
72     hist = hist[:,1:lastiter]
73 end
74
75 if normg > tol
76     @warn "Did not converge"
77 end
78
79 return x,f,g,hist,alpha_list,histx
80 end
81
Out[195]:

```

newton (generic function with 1 method)

4. Prepare a plot of the step lengths as in part 2.

Firstly, we compute the function starting from the point (1.2, 1.2)

In [196]:

```

1 x,fx,gx,hist,alpha_list3,histx3= newton(fgh, [1.2,1.2];
2 maxiter=16000, tol=1.0e-8,quiet=false,
3 histx=histx, hista=hista, gamma=0.01);

```

iter	val	normg	fdiff
1	5.80e+00	1.16e+02	Inf
2	3.84e-02	4.00e-01	5.76e+00
3	1.88e-02	4.39e+00	1.96e-02
4	4.29e-03	6.15e-01	1.45e-02
5	9.03e-04	1.14e+00	3.39e-03
6	1.85e-05	3.25e-02	8.85e-04
7	3.40e-08	7.19e-03	1.85e-05
8	3.23e-14	1.36e-06	3.40e-08
9	1.09e-25	1.29e-11	3.23e-14


In this case, it converges.

Let's plot the values of α

In [197]:

```
1 plot(alpha_list3)
```

Out[197]:



This plot is much more clear than the ones from the gradient descent.

Let's see the contour plot.

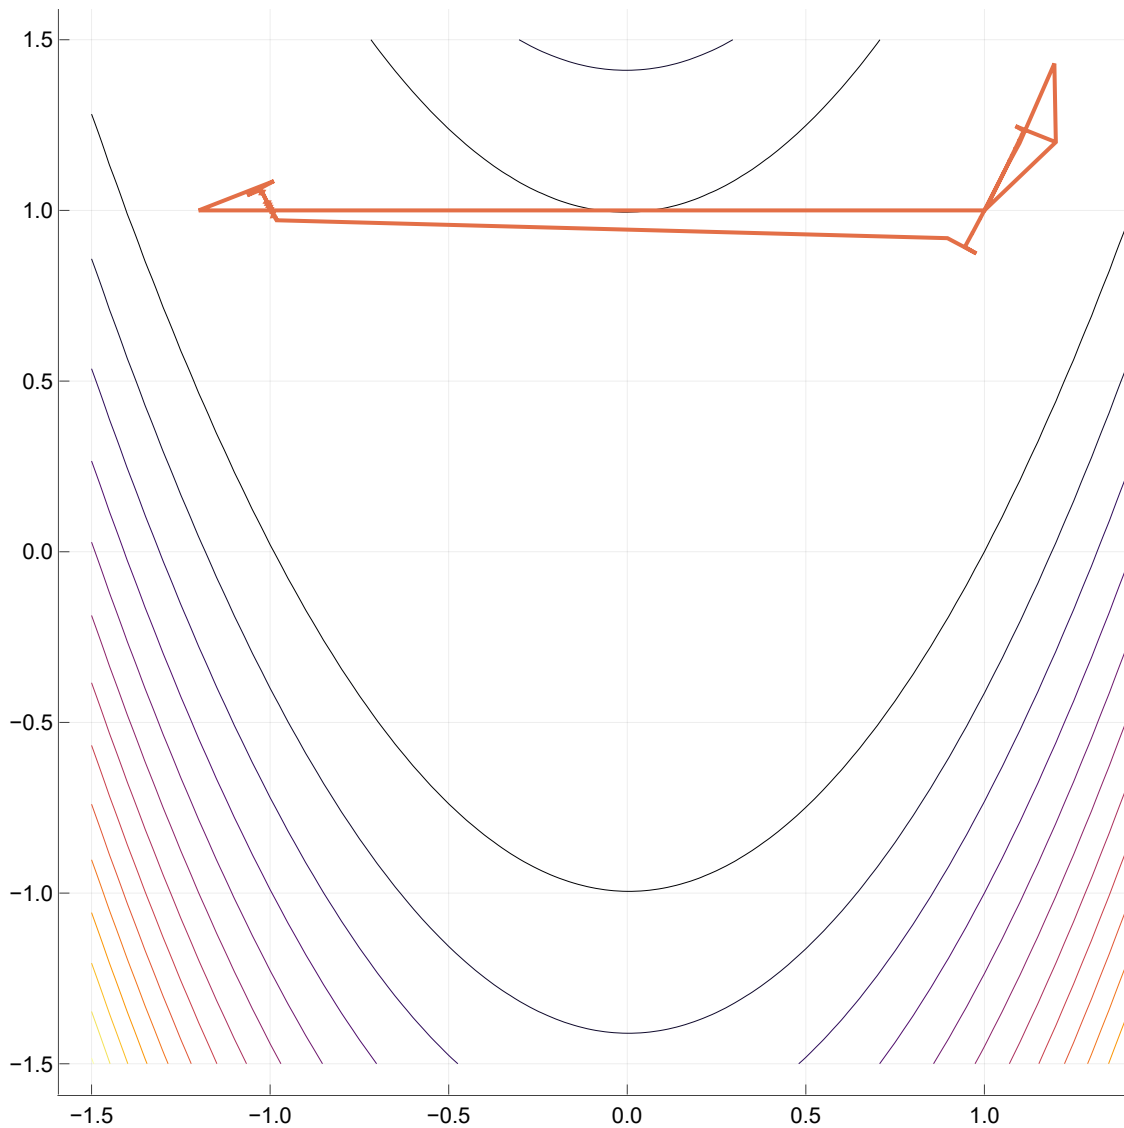
In [198]:

```

1 ezcontour(-1.5:0.05:1.5, -1.5:0.05:1.5,
2 MultivariateProblems.UnconstrainedProblems.examples["Rosenbrock"].f)
3 #histx1 = histx1[1:3]
4 plot!(map(first,histx3),map(x->x[2], histx3), linewidth=2)

```

Out[198]:



Now, we compute the function starting from the point $(-1.2, 1)$

In [199]:

```

1 x,fx,gx,hist,alpha_list4,histx4= newton(fgh, [-1.2,1];
2 maxiter=16000, tol=1.0e-8,quiet=false,
3 histx=histx, hista=hista, gamma=0.01);

```

iter	val	normg	fdiff
1	2.42e+01	2.16e+02	Inf
2	4.73e+00	4.64e+00	1.95e+01
3	4.09e+00	2.60e+01	6.44e-01
4	3.23e+00	1.06e+01	8.59e-01
5	3.21e+00	2.21e+01	1.48e-02
6	1.94e+00	3.49e+00	1.27e+00
7	1.60e+00	7.42e+00	3.42e-01
8	1.18e+00	4.13e+00	4.22e-01
9	9.22e-01	8.63e+00	2.56e-01
10	5.97e-01	1.59e+00	3.25e-01
11	4.53e-01	5.21e+00	1.45e-01
12	2.81e-01	1.99e+00	1.72e-01
13	2.11e-01	7.20e+00	6.94e-02
14	8.90e-02	4.84e-01	1.22e-01
15	5.15e-02	3.22e+00	3.75e-02
16	2.00e-02	1.00e+00	3.15e-02
17	7.17e-03	2.21e+00	1.28e-02
18	1.07e-03	1.96e-01	6.10e-03
19	7.78e-05	3.10e-01	9.92e-04
20	2.82e-07	3.23e-03	7.75e-05
21	8.52e-12	1.06e-04	2.82e-07
22	3.74e-21	3.73e-10	8.52e-12


In this case, the method also converges.

Let's see the values of α

In [200]:

```
1 plot(alpha_list4)
```

Out[200]:



Now, let's see the contour plot


In [201]:

```

1 ezcontour(-1.5:0.05:1.5, -1.5:0.05:1.5,
2 MultivariateProblems.UnconstrainedProblems.examples["Rosenbrock"].f)
3 #histx1 = histx1[1:3]
4 plot!(map(first,histx4),map(x->x[2], histx4), linewidth=2)

```

Out[201]:



5. Discuss any notable differences or similarities.

The main difference between the Gradient Descent and the Newton's Method is the number of iterations. The Newton's method needs less iterations.

Another difference that can be seen from the plots is that the value of α is close to 0 in most of the iterations when using the Gradient descent. While the value of α is close to 1 in most of the iterations when using the Newton's Method.

PROBLEM 2

(Exercise 4.2 in Kochenderfer and Wheeler) The first Wolfe condition requires

$f(x + \alpha p) \leq f(x) + c_1 \alpha p^T g$. What is the maximum step, length α that satisfies this condition given

that $f(x) = 5 + x_1^2 + x_2^2$, $x = [-1, 1]^T$, $p = [1, 0]^T$, $c_1 = 10^{-4}$

With the information of the statement, we have to find the maximum α that satisfies the first Wolfe condition.

Firstly, let's compute g

$$g = \frac{\partial f(x)}{\partial x_1} i + \frac{\partial f(x)}{\partial x_2} j = 2x_1 i + 2x_2 j = \begin{bmatrix} 2x_1 & 2x_2 \end{bmatrix}$$

$$\text{As } x = [-1, 1]^T, g = \begin{bmatrix} -2 & 2 \end{bmatrix}$$

Let's compute $f(x + \alpha p)$

$$f(x + \alpha p) = 5 + (x_1 + \alpha p_1)^2 + (x_2 + \alpha p_2)^2$$

$$\text{As } x = [-1 \ 1] \text{ and } p = [1 \ 0],$$

$$f(x + \alpha p) = 5 + (-1 + \alpha \cdot 1)^2 + (1 + \alpha \cdot 0)^2 = 5 + 1 + \alpha^2 - 2\alpha + 1 = \alpha^2 - 2\alpha + 7$$

Let's compute $f(x)$

$$\text{As } x = [-1 \ 1],$$

$$f(x) = 5 + (-1)^2 + (1)^2 = 7$$

Let's compute $p^T g$

$$p^T g = [1 \ 0] \cdot \begin{bmatrix} -2 \\ 2 \end{bmatrix} = -2$$

Let's substitute all the values into the first Wolfe condition

$$f(x + \alpha p) \leq f(x) + c_1 \alpha p^T g$$

$$\alpha^2 - 2\alpha + 7 \leq 7 - 2 \cdot 10^{-4} \alpha$$

$$\alpha^2 - 2\alpha + 2 \cdot 10^{-4} \alpha \leq 0$$

If we solve this inequality, we obtain

$$\alpha \leq 2(1 - 10^{-4}) \approx 1.9998$$

So, the maximum length α that satisfies this condition is: $\alpha = 1.9998$

PROBLEM 3

(Modified from Griva, Sofer, and Nash 11.5.8.) Consider the function.

$$f(x) = -q^T x + \sum_j y_j \log(C^T x)_j$$

where C is a very large matrix and computing $C^T x$ is expensive. Show how we can use the structure of the function to reduce the number of matrix vector products $C^T x$ that would be required for checking the Wolfe conditions when searching in a direction p . (Hint, compute $w = C^T p$ once, and see how to re-use it.)

The Wolfe conditions are

$$1. f(x_k + \alpha p_k) \leq f(x_k) + \alpha c_1 p_k^T g_k$$

$$2. g(x_k + \alpha p_k)^T p \geq c_2 g_k^T p_k$$

Let's compute the gradient

$$\frac{\partial f}{\partial x_i} = -q_i + \sum_j y_j \frac{c_{ij}}{(C^T x)_j}$$

When checking the wolfe conditions, we have to compute $f(x_k + \alpha p_k)$ and $g(x_k + \alpha p_k)$:

$$f(x_k + \alpha p_k) = -q^T(x_k + \alpha p_k) + \sum_j y_j \log(C^T(x_k + \alpha p_k))_j$$

$$g(x_k + \alpha p_k) = -q_i + \sum_j y_j \frac{c_{ij}}{(C^T(x_k + \alpha p_k))_j}$$

At the begining of each iteration k we can compute:

$$w = C^T p$$

$$z = C^T x_k$$

$$f(x_k)$$

$$g(x_k)$$

Every time we verify if the wolfe conditions are fulfilled we can reuse them, because $C^T x$ is linear.

When computing $f(x_k + \alpha p_k)$ and $g(x_k + \alpha p_k)$, the term $C^T(x_k + \alpha p_k)$ is the most expensive. And as we have computed some values as explained above. Now, we have:

$$C^T(x_k + \alpha p_k) = z + \alpha w$$

And this is the sum of 2 vectors

Now, the expressions $f(x_k + \alpha p_k)$ and $g(x_k + \alpha p_k)$ are:

$$f(x_k + \alpha p_k) = -q^T x_k - q^T x_k \alpha p_k + \sum_j y_j \log(z + \alpha w)_j$$

$$g(x_k + \alpha p_k) = -q_i + \sum_j y_j \frac{c_{ij}}{z_j + \alpha w_j}$$

Because we apply $C^T(x_k + \alpha p_k) = z + \alpha w$

So, by computing $w = C^T p$, we reduce the number of matrix-vector products C^T required for checking the Wolfe conditions

PROBLEM 4

Read section 3.5 in Nocedal and Wright on step-length selection. Then solve problem 3.13 (typos are mine).

1. In the notation we have from class where $L(\alpha)$ is the function projected into the current search direction, then show that the quadratic function that interpolates $L(0)$, $L'(0)$, and $L(\alpha_0)$ is given by (3.57), ie\

$$\phi_q(\alpha)_q = \left(\frac{L(\alpha_0) - L(0) - \alpha_0 L'(0)}{\alpha_0^2} \right) \alpha^2 + L'(0)\alpha + \phi(0).$$

Let's rewrite the expresion above as:

$$L_q(\alpha)_q = \left(\frac{L(\alpha_0) - L(0) - \alpha_0 L'(0)}{\alpha_0^2} \right) \alpha^2 + L'(0)\alpha + L(0).$$

We have to show that the quadratic function that interpolates $L(0)$, $L'(0)$, and $L(\alpha_0)$ is given by :

$$L_q(\alpha)_q = \left(\frac{L(\alpha_0) - L(0) - \alpha_0 L'(0)}{\alpha_0^2} \right) \alpha^2 + L'(0)\alpha + L(0).$$

Let's start with the the general form of a quadractic function:

$$L_q(\alpha) = a\alpha^2 + b\alpha + c$$

Then we have to compute the coefficients a, b and c

$$L_q(0) = a(0)^2 + b(0) + c = L(0) \implies c = L(0)$$

$$L'_q(0) = 2a(0) + b = L'(0) \implies b = L'(0)$$

$L_q(\alpha_0) = a(\alpha_0)^2 + b\alpha_0 + c = L(\alpha_0)$, As $b = L'_q(0)$ and $c = L(0)$, we substitute these values:

$$L_q(\alpha_0) = a(\alpha_0)^2 + b\alpha_0 + c = a(\alpha_0)^2 + L'_q(0)\alpha_0 + L(0) = L(\alpha_0) \text{ And we get:}$$

$$a = \frac{L(\alpha_0) - L(0) - L'_q(0)\alpha_0}{(\alpha_0)^2}$$

So, substituting the values of a, b and c , we obtain

$$L_q(\alpha)_q = \left(\frac{L(\alpha_0) - L(0) - \alpha_0 L'(0)}{\alpha_0^2} \right) \alpha^2 + L'(0)\alpha + L(0).$$

So, we obtain the same expression than the one given by (3.57).

Although er use the notation from class, the expression is the same.

2. Suppose that the sufficient decrease condition is not satisfied at α_0 . Then show that $\alpha_1 < \frac{\alpha_0}{2(1-c_1)}$

If α_0 does not satisfy the sufficient decrease conditions, then:

$$0 < L(\alpha_0) - L(0) - c_1 L'(0)\alpha_0$$

$$< L(\alpha_0) - L(0) - L'(0)\alpha_0$$

Because $L'(0) < 0$ and $c_1 < 1$. So, $a > 0$.

As L_q is convex and has a minimizer at:

$$\alpha_1 = -\frac{L'(0)\alpha_0^2}{2[L(\alpha_0) - L(0) - L'(0)\alpha_0]} \quad (3.58)$$

Notice that

$$0 < (c_1 - 1)L'(0)\alpha_0 = L(0) + c_1 L'(0)\alpha_0 - L(0) - L'(0)\alpha_0 < L(\alpha_0) - L(0) - L'(0)\alpha_0.$$

The last inequality follows from the supposition that the sufficient decrease condition is not satisfied at α_0 .

Ans by using the relations explained above, we show that:

$$\alpha_1 < -\frac{L'(0)\alpha_0^2}{2(c_1 - 1)L'(0)\alpha_0} = \frac{\alpha_0}{2(1 - c_1)}$$